

Problem statement

- we are given with image data which belongs to 16 categories...we need to build CNN model to classify the given image into its category using transfer learning with VGG16 model...

Transfer Learning

Download all the data in this [rar_file \(https://drive.google.com/open?id=1Z4TyI7FcFVEx8qdl4jO9qxvxaqLSqoEu\)](https://drive.google.com/open?id=1Z4TyI7FcFVEx8qdl4jO9qxvxaqLSqoEu) , it contains all the data required for the project. When you unrar the file you'll get the files in the following format: **path/to/the/image.tif,category**

where the categories are numbered 0 to 15, in the following order:

```
0 letter
1 form
2 email
3 handwritten
4 advertisement
5 scientific report
6 scientific publication
7 specification
8 file folder
9 news article
10 budget
11 invoice
12 presentation
13 questionnaire
14 resume
15 memo
```

There is a file named as 'labels_final.csv' , it consists of two columns. First column is path which is the required path to the images and second is the class label.

In []:

```
#the dataset that we are dealing with is quite large 3.7 GB and hence there are two methods to import the data to Colab
# Method 1- you can use gdown module to get the data directly from Google drive to Colab
# the syntax is as follows !gdown --id file_id , for ex - running the below cell will import the rvl-cdip.rar dataset
```

In []:

```
#!gdown --id 1Z4TyI7FcFVEx8qdl4jO9qxvxaqLSqoEu
```

In []:

```
# Method -2 you can also import the data using wget function
#https://www.youtube.com/watch?v=BPUfVq7RaY8
```

In []:

```
#unrar the file
#get_ipython().system_raw("unrar x rvl-cdip.rar")
```

On this image data, we are training 3 types of models as given below we have to split the data into Train and Validation data.

Model-1

1. Use [VGG-16 \(https://www.tensorflow.org/api_docs/python/tf/keras/applications/VGG16\)](https://www.tensorflow.org/api_docs/python/tf/keras/applications/VGG16) pretrained network without Fully Connected layers and initialize all the weights with Imagenet trained weights.
2. After VGG-16 network without FC layers, add a new Conv block (1 Conv layer and 1 Maxpooling), 2 FC layers and an output layer to classify 16 classes. You are free to choose any hyperparameters/parameters of conv block, FC layers, output layer.
3. Final architecture will be **INPUT --> VGG-16 without Top layers(FC) --> Conv Layer --> Maxpool Layer --> 2 FC layers --> Output Layer**
4. Print model.summary() and plot the architecture of the model.
[Reference for plotting_model \(https://www.tensorflow.org/api_docs/python/tf/keras/utils/plot_model\)](https://www.tensorflow.org/api_docs/python/tf/keras/utils/plot_model).
5. Train only new Conv block, FC layers, output layer. Don't train the VGG-16 network.

In []:

```
!curl --header "Host: www.kaggle.com" --header "User-Agent: Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML,
```

% Total	% Received	% Xferd	Average Speed	Time	Time	Time	Current
			Dload Upload	Total	Spent	Left	Speed
0	0	0	0	0	---	---	0
100	4442M	100	4442M	0	0	20.7M	0 0:03:33 0:03:33 ---:---: 22.6M

In []:

```
#!wget --header='Host: www.kaggle.com' --header='User-Agent: Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML,
```

In []:

```
#!kaggle datasets download -d brahma0545/aaic-assignment-tl
```

In []:

```
import zipfile

with zipfile.ZipFile('/content/archive.zip', 'r') as zip_ref:
    zip_ref.extractall()
```

In []:

```
import matplotlib.pyplot as plt # importing the libraries
import pandas as pd
import numpy as np
import seaborn as sns
import tensorflow as tf
import datetime, os
from tensorflow import keras
from keras.models import Model
```

In []:

```
#get_ipython().system_raw("unrar x archive.zip") # extracting the uploaded file
```

In []:

```
df=pd.read_csv("labels_final.csv")
```

In []:

```
df.head(5)
```

Out[6]:

	path	label
0	imagesv/v/o/h/voh71d00/509132755+-2755.tif	3
1	imagesl/l/x/t/xt19d00/502213303.tif	3
2	imagesx/x/e/d/xed05a00/2075325674.tif	2
3	imageso/o/j/b/ojb60d00/517511301+-1301.tif	3
4	imagesq/q/z/k/qzk17e00/2031320195.tif	7

In []:

```
labels_dict={ 0 : "letter",
              1 : "form",
              2 : "email",
              3 : "handwritten",
              4 : "advertisement",
              5 : "scientific report",
              6 : "scientific publication",
              7 : "specification",
              8 : "file folder",
              9 : "news article",
              10 : " budget",
              11 : "invoice",
              12 : " presentation",
              13 : "questionnaire",
              14 : "resume",
              15 : "memo"}
```

In []:

```
df['label']=df['label'].apply(lambda x:labels_dict[x])
df.head(5)
```

Out[6]:

	path	label
0	imagesv/v/o/h/voh71d00/509132755+-2755.tif	handwritten
1	imagesl/l/x/t/lxt19d00/502213303.tif	handwritten
2	imagesx/x/e/d/xed05a00/2075325674.tif	email
3	imageso/o/j/b/ojb60d00/517511301+-1301.tif	handwritten
4	imagesq/q/z/k/qzk17e00/2031320195.tif	specification

In []:

```
from keras_preprocessing.image import ImageDataGenerator
datagen = ImageDataGenerator(rescale=1/255., validation_split=0.2) #image generator

print("-----TRAIN DATA-----") # train data
train_generator = datagen.flow_from_dataframe(dataframe=df, directory="/content/data_final",
                                              x_col='path',
                                              y_col='label', # using flow from data frame
                                              target_size=(256,256),
                                              class_mode='categorical',
                                              batch_size=32,
                                              drop_remainder=True,
                                              subset='training',
                                              seed=7)

print("-----CROSS VALIDATION DATA-----") # cross validation data

validation_generator = datagen.flow_from_dataframe(dataframe=df, directory="/content/data_final",
                                                  x_col='path',
                                                  y_col='label',
                                                  target_size=(256,256),
                                                  class_mode='categorical',
                                                  batch_size=32,
                                                  drop_remainder=True,
                                                  subset='validation',
                                                  seed=7)
```

```
-----TRAIN DATA-----
Found 38400 validated image filenames belonging to 16 classes.
-----CROSS VALIDATION DATA-----
Found 9600 validated image filenames belonging to 16 classes.
```

In []:

```
import tensorflow as tf
from tensorflow.keras.layers import Input, Lambda, Dense, Flatten
from tensorflow.keras.models import Model
from tensorflow.keras.applications.vgg16 import VGG16
from tensorflow.keras.applications.vgg16 import preprocess_input
from tensorflow.keras.preprocessing import image
from tensorflow.keras.layers import Dense, Conv2D, MaxPool2D , Flatten
from tensorflow.keras.callbacks import Callback
from tensorflow.keras.callbacks import TensorBoard
```

In []:

```
%load_ext tensorboard
```

In []:

```
logdir="logs/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S") # tensorboard
tensorboard_callback = TensorBoard(log_dir=logdir, histogram_freq=0)
```

In []:

```
IMAGE_SIZE = [256, 256] #pre trained vgg16 model
model = VGG16(input_shape=IMAGE_SIZE + [3], weights='imagenet', include_top=False)
```

```
model.summary() #pre trained vgg16 model
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5 (https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5)

58889256/58889256 [=====] - 4s 0us/step

Model: "vgg16"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	(None, 256, 256, 3)	0
block1_conv1 (Conv2D)	(None, 256, 256, 64)	1792
block1_conv2 (Conv2D)	(None, 256, 256, 64)	36928
block1_pool (MaxPooling2D)	(None, 128, 128, 64)	0
block2_conv1 (Conv2D)	(None, 128, 128, 128)	73856
block2_conv2 (Conv2D)	(None, 128, 128, 128)	147584
block2_pool (MaxPooling2D)	(None, 64, 64, 128)	0
block3_conv1 (Conv2D)	(None, 64, 64, 256)	295168
block3_conv2 (Conv2D)	(None, 64, 64, 256)	590080
block3_conv3 (Conv2D)	(None, 64, 64, 256)	590080
block3_pool (MaxPooling2D)	(None, 32, 32, 256)	0
block4_conv1 (Conv2D)	(None, 32, 32, 512)	1180160
block4_conv2 (Conv2D)	(None, 32, 32, 512)	2359808
block4_conv3 (Conv2D)	(None, 32, 32, 512)	2359808
block4_pool (MaxPooling2D)	(None, 16, 16, 512)	0
block5_conv1 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv2 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv3 (Conv2D)	(None, 16, 16, 512)	2359808
block5_pool (MaxPooling2D)	(None, 8, 8, 512)	0
=====		
Total params: 14,714,688		
Trainable params: 14,714,688		
Non-trainable params: 0		

In []:

```
train_steps = train_generator.n//train_generator.batch_size
validation_steps = validation_generator.n//validation_generator.batch_size
```

In []:

```
#model_1
for layer in model.layers:
    layer.trainable = False
#Adding custom Layers
x = model.output
x = Conv2D(filters=512,kernel_size=(3,3),padding="same", activation="relu")(x)
x = MaxPool2D(2,2)(x)
x = Flatten()(x)
x = Dense(256, activation="relu")(x)
x = Dense(128, activation="relu")(x)
output = Dense(16, activation="softmax")(x)
# creating the final model
model_1 = Model(inputs = model.input, outputs = output)
# compile the model
model_1.compile(loss = "categorical_crossentropy", optimizer = 'Adam', metrics=["accuracy"])
```

In []:

```
# summary of the model_1
model_1.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 256, 256, 3)]	0
block1_conv1 (Conv2D)	(None, 256, 256, 64)	1792
block1_conv2 (Conv2D)	(None, 256, 256, 64)	36928
block1_pool (MaxPooling2D)	(None, 128, 128, 64)	0
block2_conv1 (Conv2D)	(None, 128, 128, 128)	73856
block2_conv2 (Conv2D)	(None, 128, 128, 128)	147584
block2_pool (MaxPooling2D)	(None, 64, 64, 128)	0
block3_conv1 (Conv2D)	(None, 64, 64, 256)	295168
block3_conv2 (Conv2D)	(None, 64, 64, 256)	590080
block3_conv3 (Conv2D)	(None, 64, 64, 256)	590080
block3_pool (MaxPooling2D)	(None, 32, 32, 256)	0
block4_conv1 (Conv2D)	(None, 32, 32, 512)	1180160
block4_conv2 (Conv2D)	(None, 32, 32, 512)	2359808
block4_conv3 (Conv2D)	(None, 32, 32, 512)	2359808
block4_pool (MaxPooling2D)	(None, 16, 16, 512)	0
block5_conv1 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv2 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv3 (Conv2D)	(None, 16, 16, 512)	2359808
block5_pool (MaxPooling2D)	(None, 8, 8, 512)	0
conv2d (Conv2D)	(None, 8, 8, 512)	2359808
max_pooling2d (MaxPooling2D)	(None, 4, 4, 512)	0
flatten (Flatten)	(None, 8192)	0
dense (Dense)	(None, 256)	2097408
dense_1 (Dense)	(None, 128)	32896
dense_2 (Dense)	(None, 16)	2064
=====		
Total params: 19,206,864		
Trainable params: 4,492,176		
Non-trainable params: 14,714,688		

In []:

```
#fitting the model_1
model_1.fit_generator(train_generator,steps_per_epoch=train_steps, epochs=5,
                    validation_data=validation_generator,validation_steps=validation_steps,callbacks=[tensorboard_callback])
```

<ipython-input-19-f33c75d26d90>:4: UserWarning: `Model.fit_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.

```
model_1.fit_generator(train_generator,steps_per_epoch=train_steps, epochs=5,

Epoch 1/5
1200/1200 [=====] - 333s 266ms/step - loss: 1.3061 - accuracy: 0.5939 - val_loss: 1.0037 -
val_accuracy: 0.6965
Epoch 2/5
1200/1200 [=====] - 352s 294ms/step - loss: 0.8966 - accuracy: 0.7256 - val_loss: 0.8949 -
val_accuracy: 0.7352
Epoch 3/5
1200/1200 [=====] - 315s 262ms/step - loss: 0.7528 - accuracy: 0.7684 - val_loss: 0.8907 -
val_accuracy: 0.7389
Epoch 4/5
1200/1200 [=====] - 309s 258ms/step - loss: 0.6509 - accuracy: 0.7988 - val_loss: 0.8663 -
val_accuracy: 0.7506
Epoch 5/5
1200/1200 [=====] - 314s 262ms/step - loss: 0.5694 - accuracy: 0.8202 - val_loss: 0.8897 -
val_accuracy: 0.7464
```

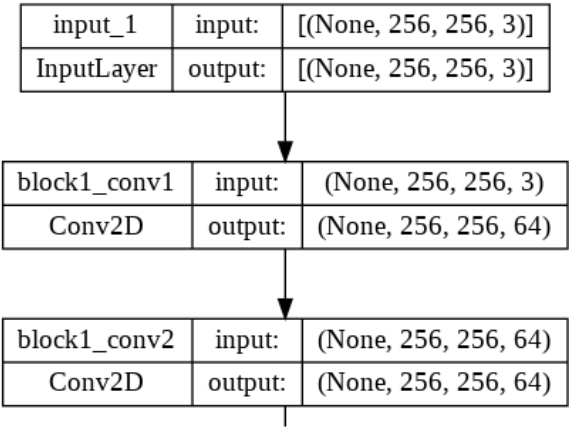
Out[19]:

<keras.callbacks.History at 0x7f37f54c7220>

In []:

```
# model graphs
tf.keras.utils.plot_model(
    model_1, to_file='model_1.png', show_shapes=True, show_layer_names=True,
    rankdir='TB', expand_nested=False, dpi=96
)
```

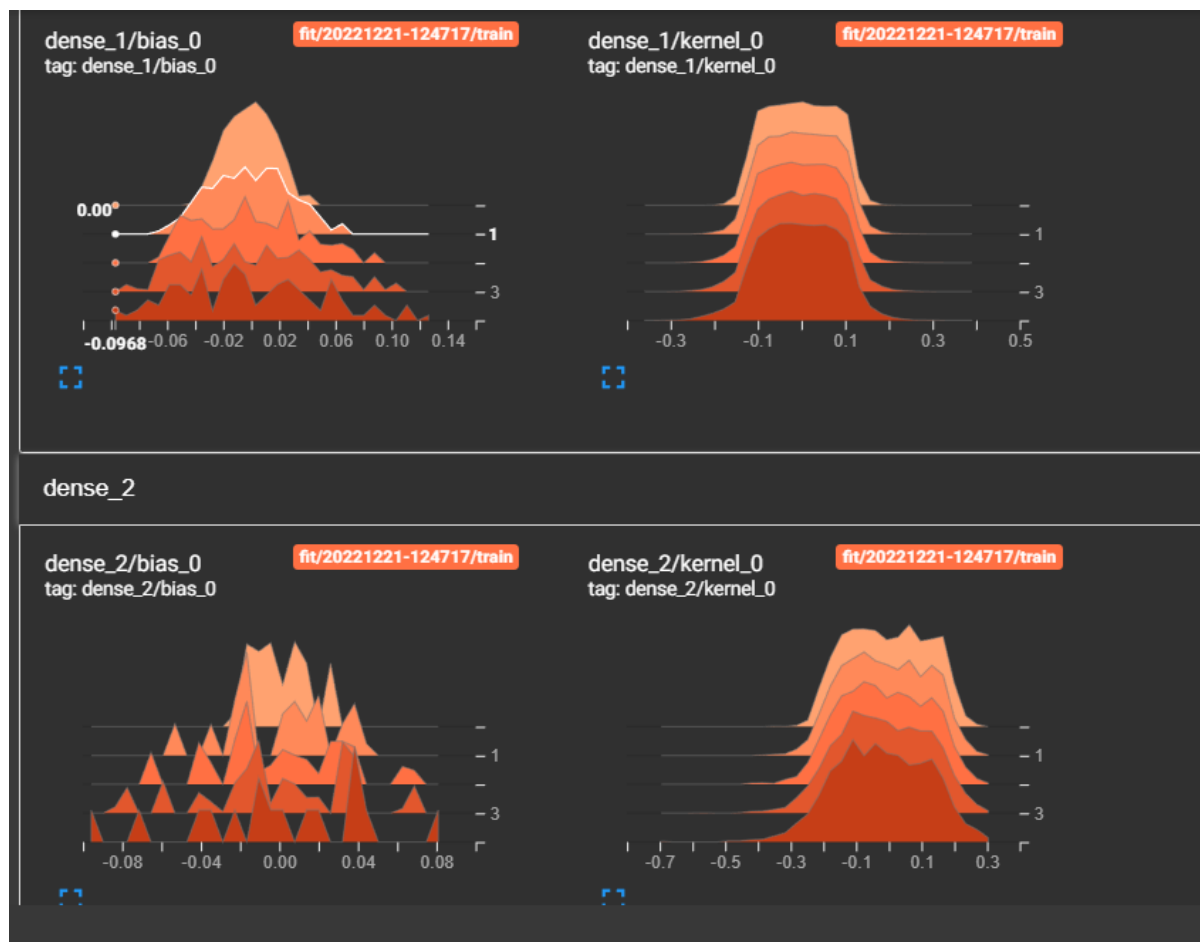
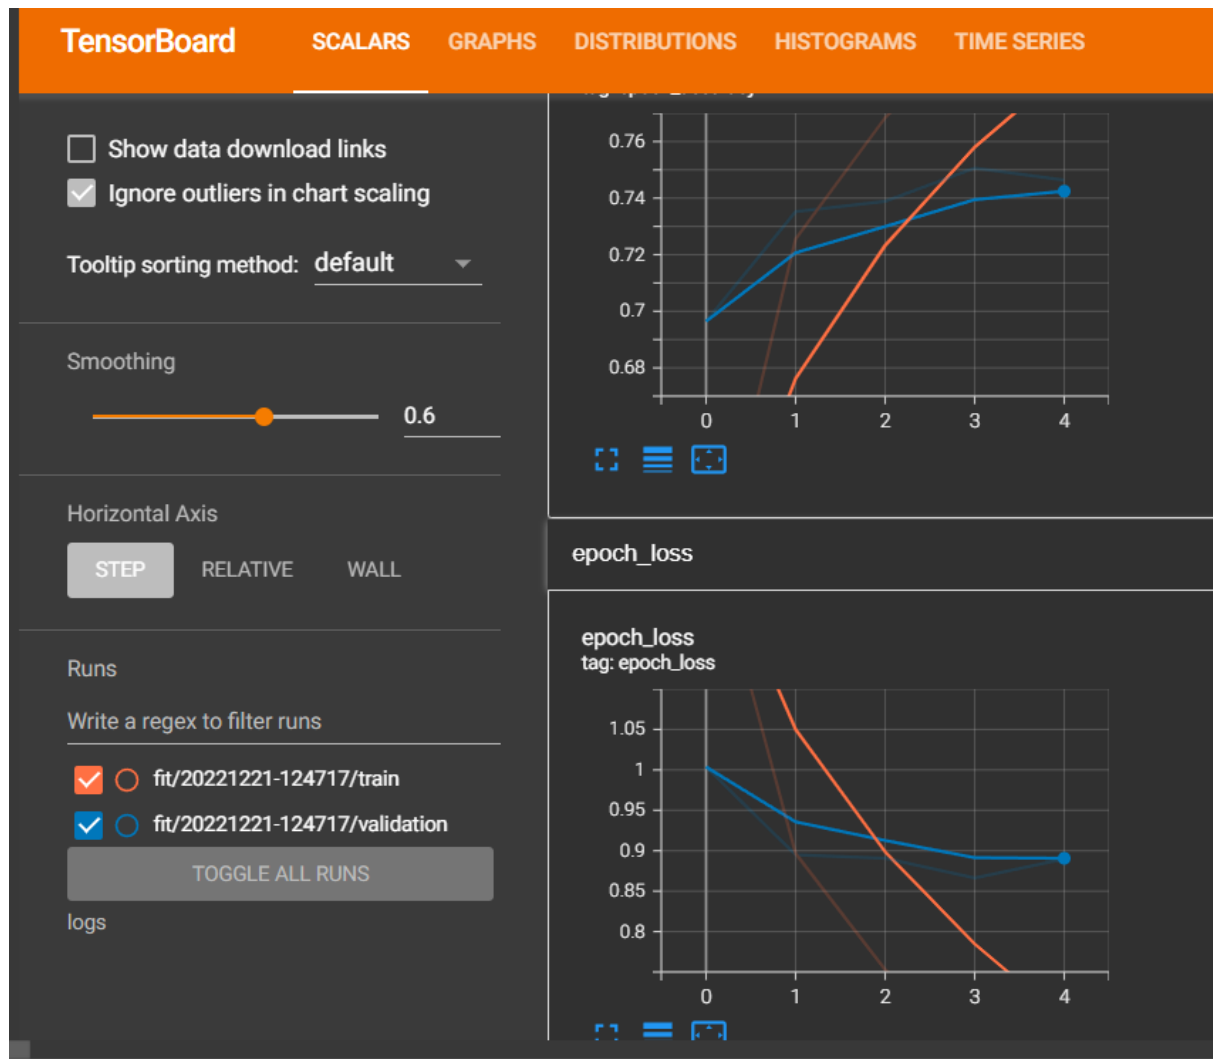
Out[20]:

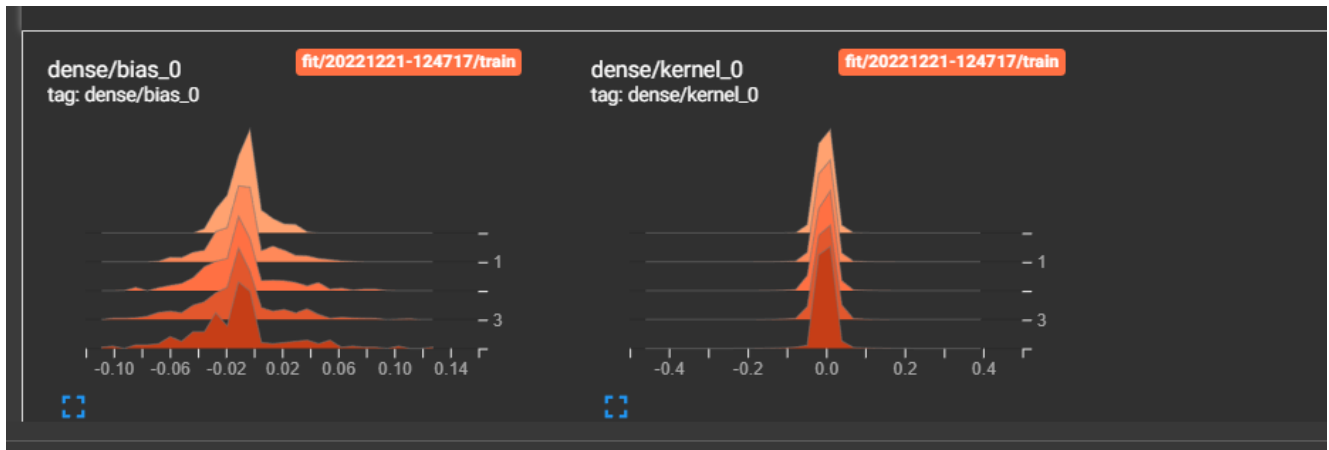


In []:

```
!rm -rf ./logs/
```

In []:



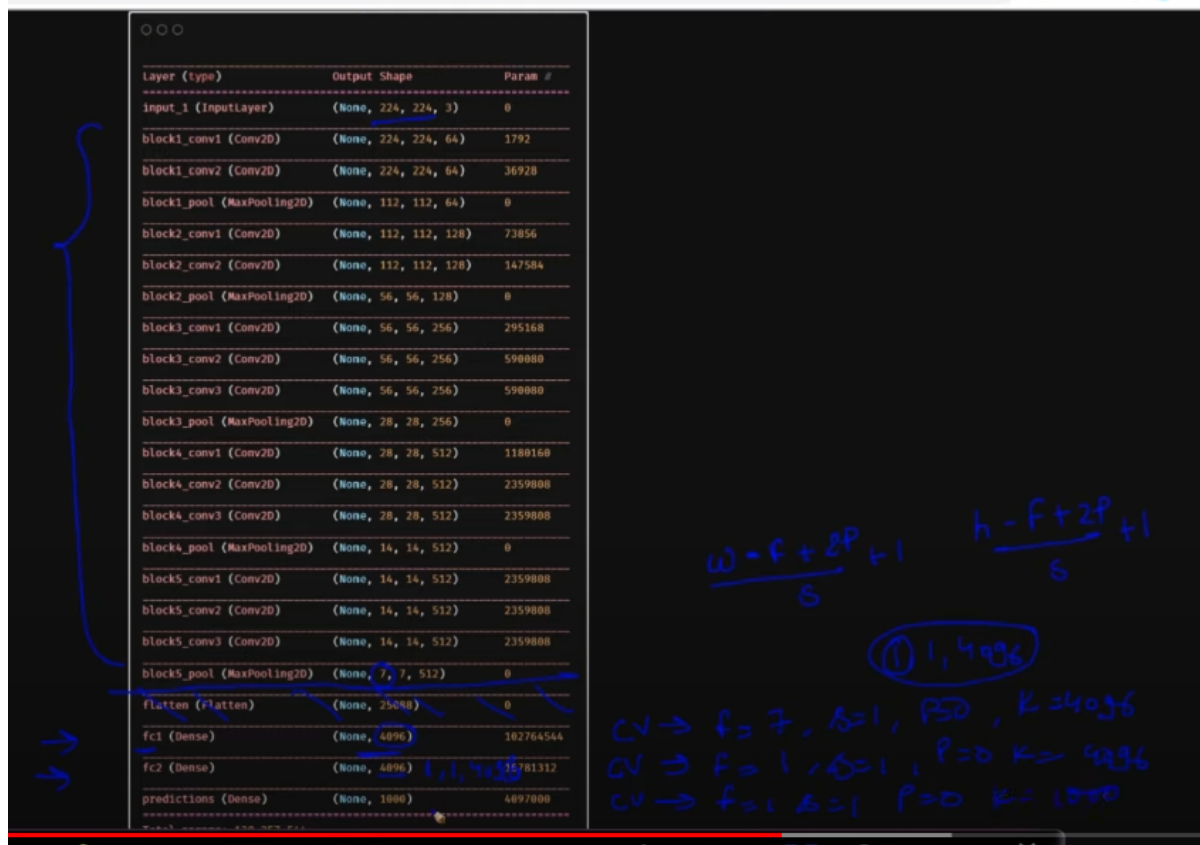


Observation

- we can observe that 1/4th (which is 4,492,176 / 19,206,864) of the parameters are trainable
- after 5 epoch we got the train accuracy of 82% and val_accuracy of 74.6%
- the weights in the last 3rd and 2nd dense layers are well distributed and the last layers weights have many values closer to zero since many of the features are zeroed out from the 128 dense neurons layer to 16 neurons layer for classification.

Model-2

1. Use [VGG-16](https://www.tensorflow.org/api_docs/python/tf/keras/applications/VGG16) (https://www.tensorflow.org/api_docs/python/tf/keras/applications/VGG16) pretrained network without Fully Connected layers and initialize all the weights with Imagenet trained weights.
2. After VGG-16 network without FC layers, don't use FC layers, use conv layers only as Fully connected layer. Any FC layer can be converted to a CONV layer. This conversion will reduce the No of Trainable parameters in FC layers. For example, an FC layer with K=4096 that is looking at some input volume of size 7x7x512 can be equivalently expressed as a CONV layer with F=7, P=0, S=1, K=4096. In other words, we are setting the filter size to be exactly the size of the input volume, and hence the output will simply be 1x1x4096 since only a single depth column "fits" across the input volume, giving identical result as the initial FC layer. You can refer [this](http://cs231n.github.io/convolutional-networks/#convert) (<http://cs231n.github.io/convolutional-networks/#convert>) link to better understanding of using Conv layer in place of fully connected layers.
3. Final architecture will be VGG-16 without FC layers (without top), 2 Conv layers identical to FC layers, 1 output layer for 16 class classification. **INPUT --> VGG-16 without Top layers(FC) --> 2 Conv Layers identical to FC --> Output Layer**
4. Print model.summary() and plot the architecture of the model. [Reference for plotting model](https://www.tensorflow.org/api_docs/python/tf/keras/utils/plot_model) (https://www.tensorflow.org/api_docs/python/tf/keras/utils/plot_model).
5. Train only last 2 Conv layers identical to FC layers, 1 output layer. Don't train the VGG-16 network.



Layer (type)	Output Shape	Param #
input_1 (InputLayer)	(None, 224, 224, 3)	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590880
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590880
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
flatten (Flatten)	(None, 25088)	0
fc1 (Dense)	(None, 4096)	102764544
fc2 (Dense)	(None, 4096)	16781312
predictions (Dense)	(None, 1000)	4097000

Handwritten notes on the right side of the table:

$$W = \frac{f+2P}{s} + 1$$

$$h = \frac{f+2P}{s} + 1$$

① 1,4096

CV $\rightarrow f=7, s=1, P=0, K=4096$
 CV $\rightarrow f=1, s=1, P=0, K=4096$
 CV $\rightarrow f=1, s=1, P=0, K=1000$

In []:

```
#model_2
for layer in model.layers:
    layer.trainable = False
#Adding custom Layers
x = model.output
x = Conv2D(filters=4096, kernel_size=8, strides=1, activation="relu")(x)
x = Conv2D(filters=4096, kernel_size=1, strides=1, activation="relu")(x)
x = Flatten()(x)
# creating the final model
output = Dense(16, activation="softmax")(x)
model_2 = Model(inputs = model.input, outputs = output)
# compile the model
model_2.compile(loss="categorical_crossentropy", optimizer = 'Adam', metrics=['accuracy'])
```

In []:

```
# summary of the model_2
model_2.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 256, 256, 3)]	0
block1_conv1 (Conv2D)	(None, 256, 256, 64)	1792
block1_conv2 (Conv2D)	(None, 256, 256, 64)	36928
block1_pool (MaxPooling2D)	(None, 128, 128, 64)	0
block2_conv1 (Conv2D)	(None, 128, 128, 128)	73856
block2_conv2 (Conv2D)	(None, 128, 128, 128)	147584
block2_pool (MaxPooling2D)	(None, 64, 64, 128)	0
block3_conv1 (Conv2D)	(None, 64, 64, 256)	295168
block3_conv2 (Conv2D)	(None, 64, 64, 256)	590080
block3_conv3 (Conv2D)	(None, 64, 64, 256)	590080
block3_pool (MaxPooling2D)	(None, 32, 32, 256)	0
block4_conv1 (Conv2D)	(None, 32, 32, 512)	1180160
block4_conv2 (Conv2D)	(None, 32, 32, 512)	2359808
block4_conv3 (Conv2D)	(None, 32, 32, 512)	2359808
block4_pool (MaxPooling2D)	(None, 16, 16, 512)	0
block5_conv1 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv2 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv3 (Conv2D)	(None, 16, 16, 512)	2359808
block5_pool (MaxPooling2D)	(None, 8, 8, 512)	0
conv2d (Conv2D)	(None, 1, 1, 4096)	134221824
conv2d_1 (Conv2D)	(None, 1, 1, 4096)	16781312
flatten (Flatten)	(None, 4096)	0
dense (Dense)	(None, 16)	65552
Total params: 165,783,376		
Trainable params: 151,068,688		
Non-trainable params: 14,714,688		

In []:

```
#fitting model_2
model_2.fit_generator(train_generator, steps_per_epoch=train_steps, epochs=5, verbose=1, callbacks=[tensorboard_callback])
```

<ipython-input-20-8d58a9e47380>:2: UserWarning: `Model.fit_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.

model_2.fit_generator(train_generator, steps_per_epoch=train_steps, epochs=5, verbose=1, callbacks=[tensorboard_callback])

Epoch 1/5

1200/1200 [=====] - 472s 394ms/step - loss: 0.6773 - accuracy: 0.7921

Epoch 2/5

1200/1200 [=====] - 472s 394ms/step - loss: 0.6032 - accuracy: 0.8122

Epoch 3/5

1200/1200 [=====] - 472s 393ms/step - loss: 0.5321 - accuracy: 0.8344

Epoch 4/5

1200/1200 [=====] - 474s 395ms/step - loss: 0.4828 - accuracy: 0.8484

Epoch 5/5

1200/1200 [=====] - 473s 394ms/step - loss: 0.4370 - accuracy: 0.8649


Out[20]:

<keras.callbacks.History at 0x7fae405bcdc0>

In []:

```
# model graphs
tf.keras.utils.plot_model(
    model_2, to_file='model_2.png', show_shapes=False, show_layer_names=True,
    rankdir='TB', expand_nested=False, dpi=96
)
```

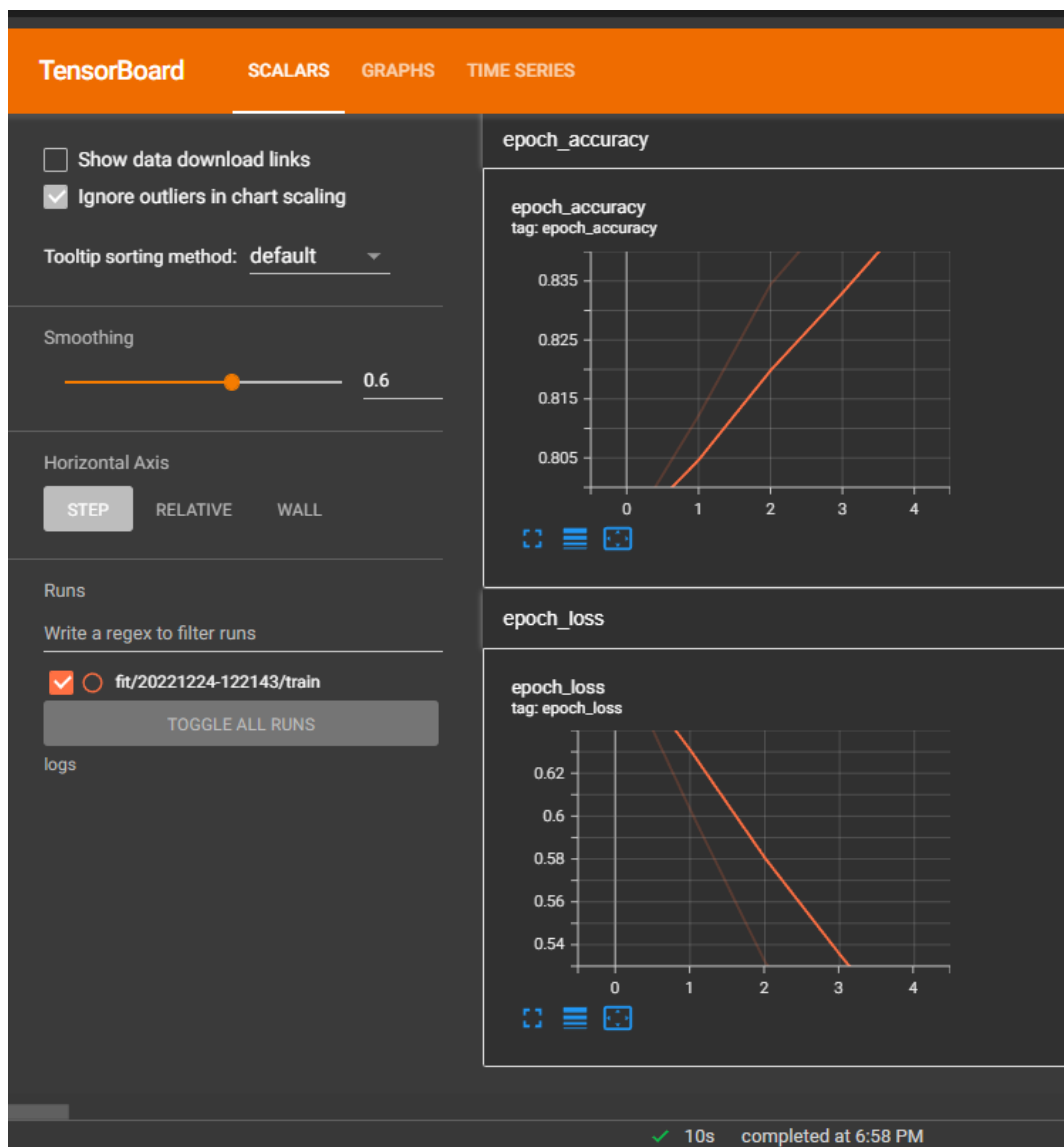
Out[21]:



In []:

```
!rm -rf ./logs/
```

In []:



Observation

- we can observe that 911/1000th (which is 151,068,688/ 165,783,376) of the parameters are trainable
- after 5 epoch we got the train accuracy of 86.49%

Model-3

1. Use same network as Model-2 'INPUT --> VGG-16 without Top layers(FC) --> 2 Conv Layers identical to FC --> Output Layer' and train only Last 6 Layers of VGG-16 network, 2 Conv layers identical to FC layers, 1 output layer.

In []:

```
#model_1
for layer in model.layers:
    layer.trainable = False
for layer in model.layers[-6:]: # training last 6 layers of vgg16
    layer.trainable = True
    print("Layer '%s' is trainable" % layer.name)
```

```
Layer 'block4_conv3' is trainable
Layer 'block4_pool' is trainable
Layer 'block5_conv1' is trainable
Layer 'block5_conv2' is trainable
Layer 'block5_conv3' is trainable
Layer 'block5_pool' is trainable
```

In []:

```
#model_3
#Adding custom Layers
x = model.output
x = Conv2D(filters=4096, kernel_size=8, strides=1, activation="relu")(x)
x = Conv2D(filters=4096, kernel_size=1, strides=1, activation="relu")(x)
x = Flatten()(x)
# creating the final model
output = Dense(16, activation="softmax")(x)
model_3 = Model(inputs = model.input, outputs = output)
# compile the model
model_3.compile(loss="categorical_crossentropy", optimizer = 'Adam', metrics=['accuracy'])
```

In []:

model_3.summary()

Model: "model_2"

Layer (type)	Output Shape	Param #
input_2 (InputLayer)	[(None, 256, 256, 3)]	0
block1_conv1 (Conv2D)	(None, 256, 256, 64)	1792
block1_conv2 (Conv2D)	(None, 256, 256, 64)	36928
block1_pool (MaxPooling2D)	(None, 128, 128, 64)	0
block2_conv1 (Conv2D)	(None, 128, 128, 128)	73856
block2_conv2 (Conv2D)	(None, 128, 128, 128)	147584
block2_pool (MaxPooling2D)	(None, 64, 64, 128)	0
block3_conv1 (Conv2D)	(None, 64, 64, 256)	295168
block3_conv2 (Conv2D)	(None, 64, 64, 256)	590080
block3_conv3 (Conv2D)	(None, 64, 64, 256)	590080
block3_pool (MaxPooling2D)	(None, 32, 32, 256)	0
block4_conv1 (Conv2D)	(None, 32, 32, 512)	1180160
block4_conv2 (Conv2D)	(None, 32, 32, 512)	2359808
block4_conv3 (Conv2D)	(None, 32, 32, 512)	2359808
block4_pool (MaxPooling2D)	(None, 16, 16, 512)	0
block5_conv1 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv2 (Conv2D)	(None, 16, 16, 512)	2359808
block5_conv3 (Conv2D)	(None, 16, 16, 512)	2359808
block5_pool (MaxPooling2D)	(None, 8, 8, 512)	0
conv2d_4 (Conv2D)	(None, 1, 1, 4096)	134221824
conv2d_5 (Conv2D)	(None, 1, 1, 4096)	16781312
flatten_2 (Flatten)	(None, 4096)	0
dense_2 (Dense)	(None, 16)	65552
Total params: 165,783,376		
Trainable params: 160,507,920		
Non-trainable params: 5,275,456		

In []:

#fitting model_3

```
model_3.fit_generator(train_generator, steps_per_epoch=train_steps, epochs=2,
                    callbacks=[tensorboard_callback])
```

Epoch 1/2

<ipython-input-27-0939d2b11be3>:2: UserWarning: `Model.fit_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.

```
model_3.fit_generator(train_generator, steps_per_epoch=train_steps, epochs=2,
```

```
1200/1200 [=====] - 788s 657ms/step - loss: 2.7729 - accuracy: 0.0631
```

Epoch 2/2

```
1200/1200 [=====] - 787s 655ms/step - loss: 2.7728 - accuracy: 0.0616
```

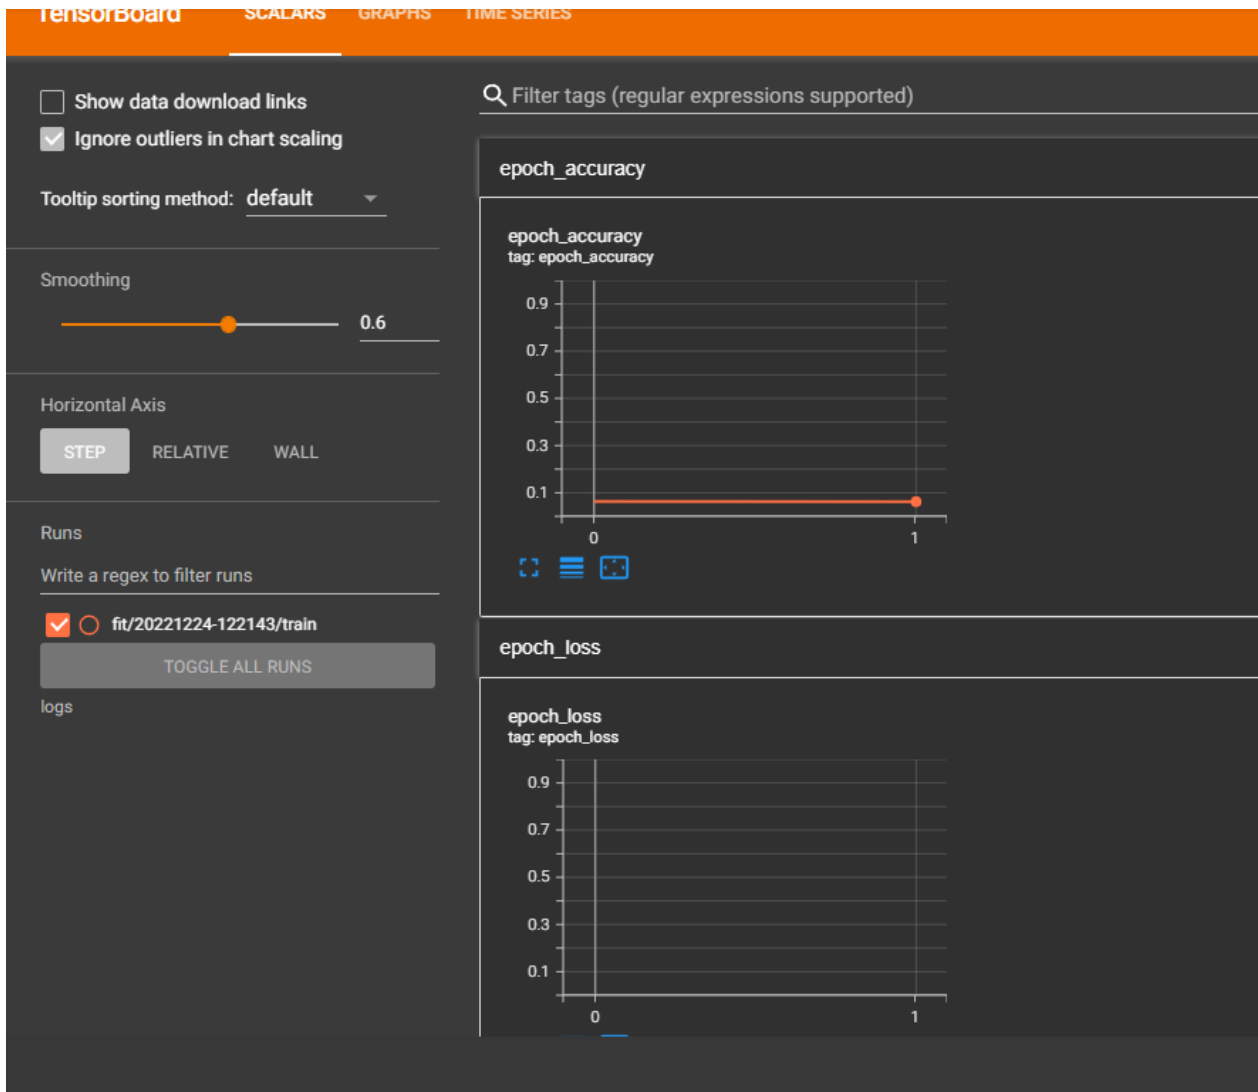

Out[27]:

```
<keras.callbacks.History at 0x7fae402a4e80>
```

In []:

```
# model graphs
tf.keras.utils.plot_model(
    model_3, to_file='model_3.png', show_shapes=False, show_layer_names=True,
    rankdir='TB', expand_nested=False, dpi=96
)
```

Out[28]:



Observation

- we can observe that 968/1000th(which is 160,507,920/ 165,783,376) of the parameters are trainable
- after 5 epoch we got the train accuracy of 6.16%.

Summary

In [6]:

```
# Please compare all your models using Prettytable Library
#http://zetcode.com/python/prettytable/
from prettytable import PrettyTable

0
x = PrettyTable()
x.field_names = ["Features", "Model", "Epochs", "Train accuracy", "CV accuracy"]
x.add_row(["Assignment 1 VGG16 w/o FC + 1 Conv + 1 MP + 2 FC + output", "CNN+Dense+softmax", "5", "82.02%", "74.64%"])
x.add_row(["Assignment 2 VGG16 w/o FC + 2 Conv + output", "CNN+softmax", "5", "86.49%", "-"])
x.add_row(["Assignment 3 VGG16 w/o FC + 2 Conv + output", "CNN+softmax", "2", "6.16%", "-"])
print(x)
```

Features	Model	Epochs	Train accuracy	CV accuracy
Assignment 1 VGG16 w/o FC + 1 Conv + 1 MP + 2 FC + output	CNN+Dense+softmax	5	82.02%	74.64%
Assignment 2 VGG16 w/o FC + 2 Conv + output	CNN+softmax	5	86.49%	-
Assignment 3 VGG16 w/o FC + 2 Conv + output	CNN+softmax	2	6.16%	-